

Processes

Linux Sysadmin Decal, Fall 2020
Ryan Chan <rrchan>

What is a process?

- A process is an instance of a program being executed
- Has a **PID** - Process ID, a numerical identifier assigned to it
- Anything that you would think of as ‘running’ is *probably* a process
 - Or in some cases, a service or a collection of processes - more on that next lecture
- Ex. Bash, vim, ping

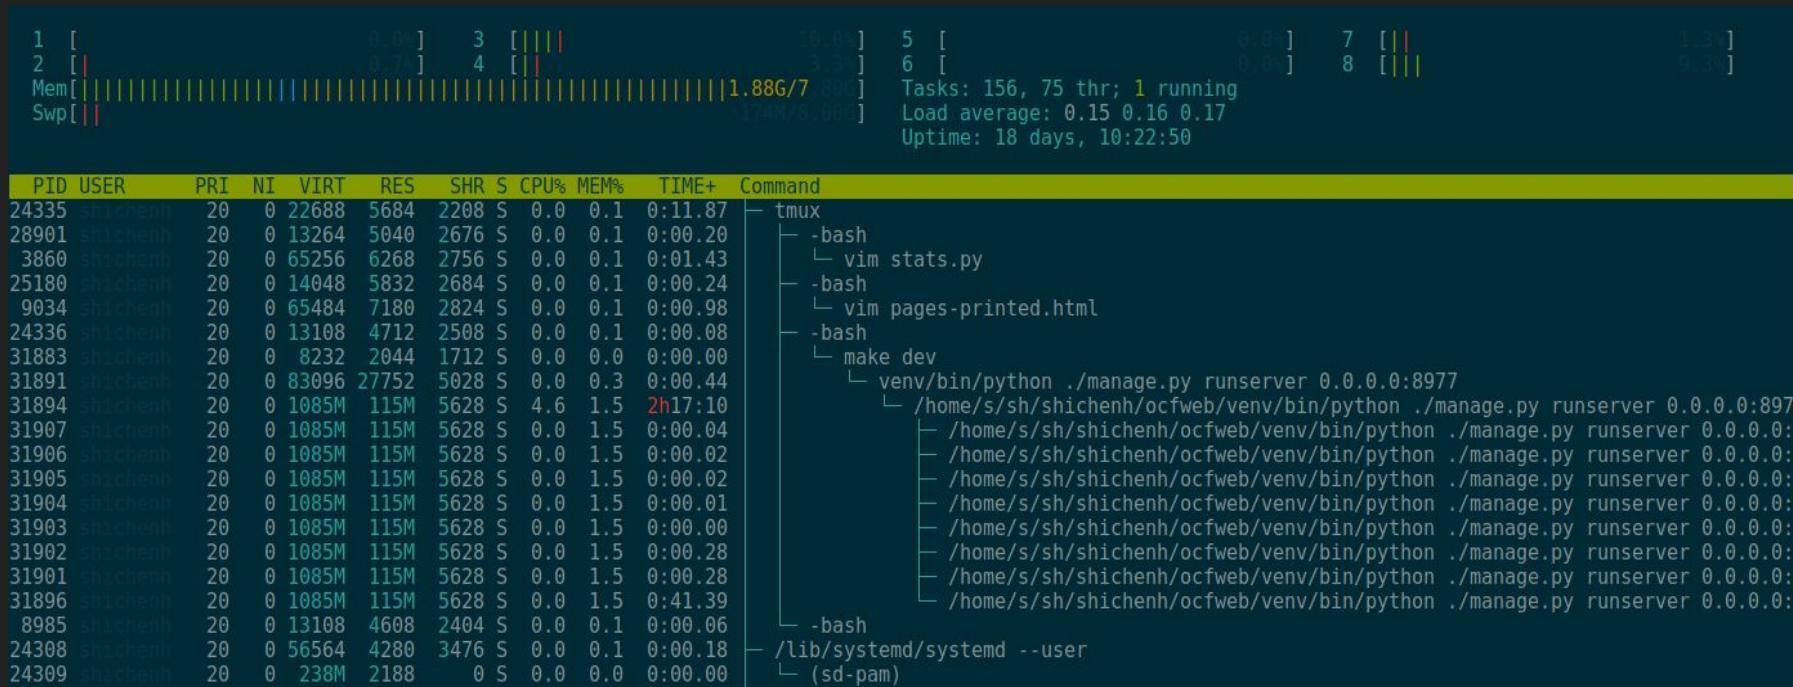
Where do processes come from?

- Processes come from other processes!
 - When a process creates another process, we call the originator the **parent** process, and the newly created process the **child** process.
 - This creates a tree-like hierarchy among processes
- Root of the process tree: **init**
 - First process run on startup, started by the kernel
 - PID 1
 - Different init systems exist. eg. You might see **systemd** as PID 1 on a system, depending on the distro you're using

Some commands

- **ps**: print running processes
 - Default: print processes in your current session
 - -e: print all processes
 - -f: print with full information
 - -H: print a tree
 - -u (--user): print a certain user's processes
- **kill, pkill**: send a signal to a process (to end it)
- **pgrep**: search for a process by its name
- **htop**: show live process stats in a nice graphical format

htop



More htop

PID	Command
1	/lib/systemd/systemd --system --deserialize 19
24854	└─ /usr/sbin/sshd -D
6392	└─ sshd: rrchan [priv]
6411	└─ sshd: rrchan@pts/0
6413	└─ -bash
7458	└─ htop
23491	└─ nginx: master process /usr/sbin/nginx -g daemon on; master_process on;
23496	└─ nginx: worker process
23494	└─ nginx: worker process
21318	└─ /lib/systemd/systemd-udevd
9077	└─ tmux
11962	└─ -bash
9118	└─ -bash
9097	└─ -bash
9078	└─ -bash

Signals: Controlling Processes

- A means of sending messages to other processes
- **SIGTERM** (15): tell a process to exit now
- **SIGINT** (2): the signal sent when you press Ctrl+C
- **SIGHUP** (1): the signal sent when the user closes the terminal window
- **SIGCHLD**: tell a process to collect child data
- **SIGKILL** (9): kill a process without giving it the chance to clean up
- **SIGSTOP** (varies; 19 on GNU/Linux): suspend a process (Ctrl + Z)
- **SIGCONT** (varies; 18 on GNU/Linux): resume a process
- The signals **SIGKILL** and **SIGSTOP** cannot be handled
- Can send using the kill, pkill, and killall commands
- e.g. kill -STOP 3276, kill -9 7283

Lifecycle of a process

1. Process is created, typically forked from a parent process
 - a. Process is assigned a PID, operating system begins tracking process info
2. Process does whatever it's supposed to be doing
3. Process ends
 - a. Operating system continues to hold onto resource info, PID, and termination status code for the process
 - b. OS sends a **SIGCHLD** signal to the parent process, informing the parent that a child process has ended and that there is info to be collected
 - c. Parent process collects the data by calling a *wait* function
 - d. OS releases all remaining data about the process

Zombie processes

- What happens if a child process ends, but the parent function does not collect the data that the OS is keeping about the child?
 - Child process will remain in the process table with its PID and other info stored, indefinitely
 - Known as a **Zombie Process**
 - Child process has ‘died’, but the parent process has not fulfilled its responsibility to ‘reap’ the child
- Not a *huge* deal, memory impact is low

Orphan Processes

- What happens to the children of a parent process if the parent process ends before its children do?
 - Child Processes are reparented, or adopted, by PID 1
 - Known as **Orphan Processes**
- What happens to processes that are both zombies and orphans?
 - Still gets sent to PID 1
 - init systems will periodically clean up zombies that are children of it

Process IO

- Every process on a Unix system initialized with three open file handles:
- Often referred to as streams
 1. **stdin** (fd 0): where a process looks for input
 2. **stdout** (fd 1): where a process can write output
 3. **stderr** (fd 2): where a process can write errors

Review: Process IO Examples

- \$ echo “stdout is the terminal”
- \$ echo “stdout is the terminal” >&1
- \$ echo “writing to stderr, also the terminal!” >&2
 - >&#: also send to the specified stream
- \$./program.sh 1>outfile 2>errfile
- \$./program.sh 2>&1
 - Send stderr to wherever stdout is going
- stdin, new way to use the cat command?
 - \$ cat < somefile.txt

Background Processes

- Noninteractive processes
- Started on the shell by appending an ‘&’ to the command
- One quirk: **stdout** and **stderr** inherited from parent shell, only **stdin** is detached
- These processes are still children of the current shell!
 - When a shell receives a **SIGHUP**, which is sent when you logout, it will also send a **SIGHUP** to each of its children, which will terminate them!

Some job control

- Jobs are child processes that a shell is taking responsibility for - not a universal Unix concept, just something that most shells do on their own
- **jobs**: See what child processes of the shell are running in the background
- **disown**: Removes a process from the shell's list of jobs
 - Will prevent the shell from sending the process a **SIGHUP**
 - The process will remain a child of the shell that ran it until the shell is closed, at which point the process will be **orphaned**
 - Gets sent to PID 1
 - Does not change the STDIN, STDOUT, or STDERR destinations of the process!
 - If one of the streams is pointing to the shell when it gets closed, the process will probably error out!

More job control

- Some more things you can do with jobs
- \$ ping -c5 ocf.berkeley.edu
- Ctrl+Z to suspend the process
- See the process listed when you run **jobs**
- **bg** to have the process resume in the background
- **fg** to bring this to foreground

Advanced process management

- Disowning background processes has its limits! What if have a process that I want running in the background all the time, and I want more tools to actually manage it?

Find out next time, in lecture b7 - Services!