

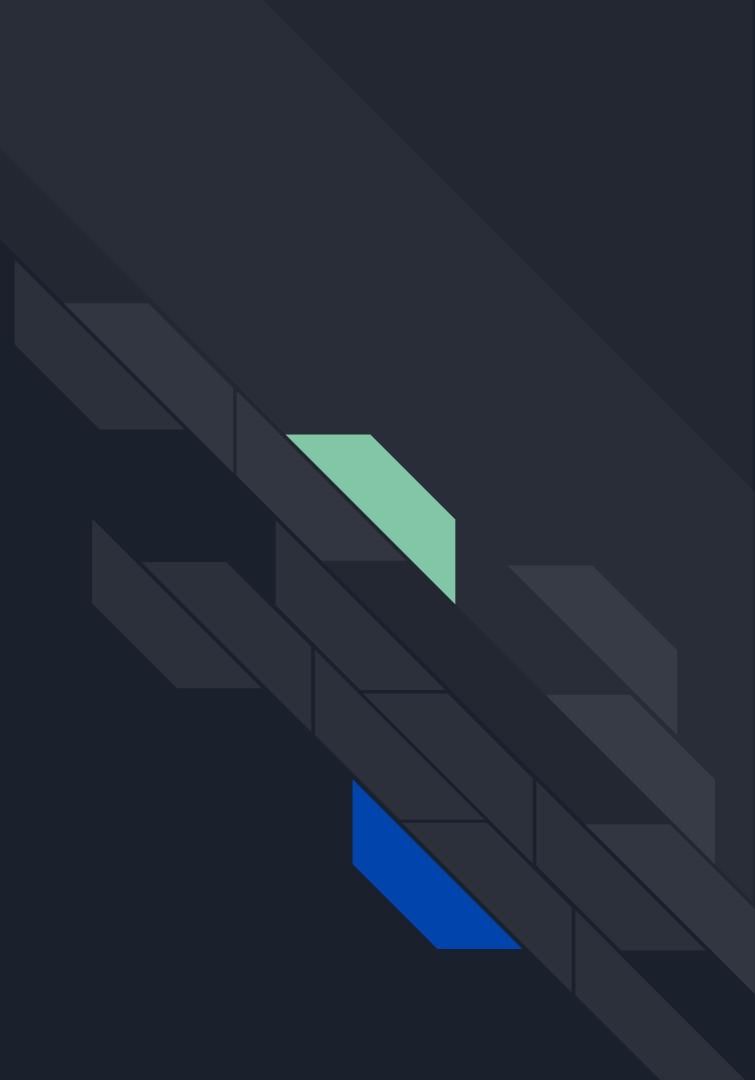


Processes and Services

Unix System Admin DeCal

Christopher Cooper <cooperc>
(based on F19 lecture by keur)

Part 1: Processes





What is a process?

A process is a single instance of a program.

Processes are isolated from one another and have their own memory, threads, etc. (Additional isolation, such as filesystem or network isolation, is also possible.)



What is a process?

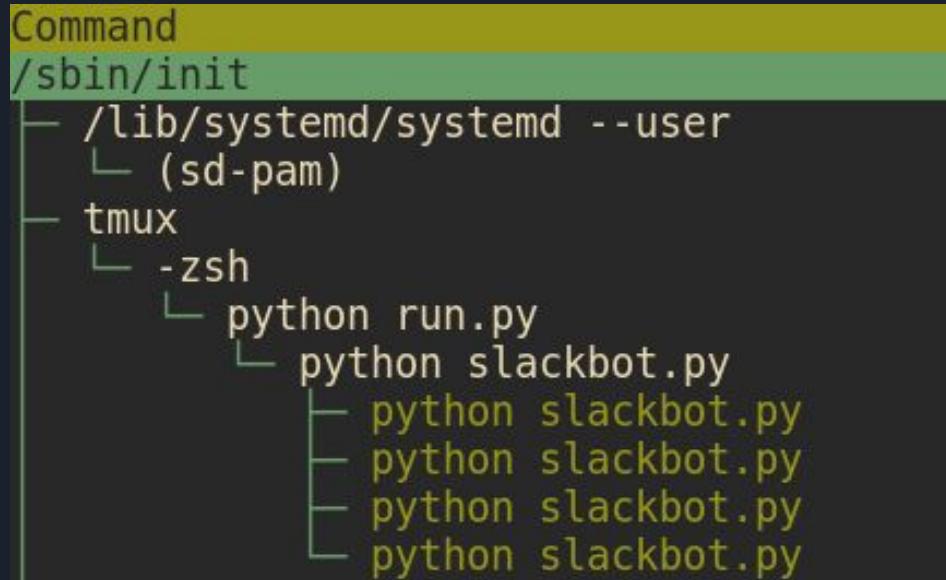
- PID: Process ID
 - PPID: Parent's PID
 - UID: User running the process
 - The program (executable) that the process is running
 - The args (command line) of the process
- (and more...)

Process Hierarchy

Each process is created by a parent

(Except PID 1)

Processes can have many children





Processes vs Threads

Both run code concurrently and can take advantage of parallelism

Threads are “lightweight” processes

All threads share virtual address space, code, static/globals, memory, resources (open files)

Processes must communicate over some interface (IPC)



Aside: Why the hell does Chrome spawn so many processes?

```
Command
└── /opt/google/chrome/chrome
    ├── /opt/google/chrome/chrome --type=gpu-process --field-trial-handle=59325916
    │   └── /opt/google/chrome/chrome --type=-broker
    ├── /opt/google/chrome/chrome-sandbox /opt/google/chrome/chrome --type=zygote
    │   └── /opt/google/chrome/chrome --type=zygote
    │       ├── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
    │       └── /opt/google/chrome/chrome --type=renderer --site-per-process --fi
```



How are processes created?

A process will **fork**(2) into two new processes, which continue from the same place

The **parent** keeps the original PID and the **child** gets a new PID

Optionally, the new process (the child) **exec**(3) and begin running a new program



How many “henlo”s get printed?

```
int main( void ) {
    fork();
    printf( "henlo: %d\n", getpid() );
}
```



How many “henlo”’s get printed?

PID 1000

```
int main( void ) {  
    fork();  
    printf( "henlo: %d\n", getpid() );  
}
```

PID 1000

```
int main( void ) {  
    fork();  
    printf( "henlo: %d\n", getpid() );  
}
```

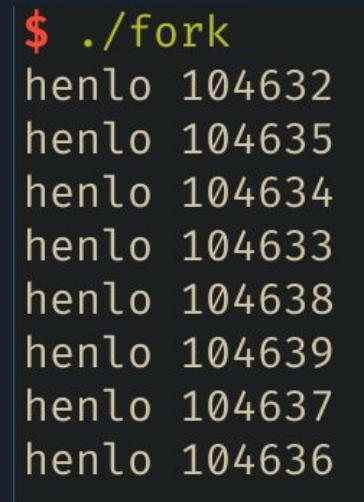
PID 2000

```
int main( void ) {  
    fork();  
    printf( "henlo: %d\n", getpid() );  
}
```



How many “henlo”s get printed?

```
int main( void ) {
    fork();    2^1
    fork();    2^2
    fork();    2^3
    printf( "henlo: %d\n", getpid() );
}
```



```
$ ./fork
henlo 104632
henlo 104635
henlo 104634
henlo 104633
henlo 104638
henlo 104639
henlo 104637
henlo 104636
```



Have you seen this code before?

```
:(){} :|:& };:
```

```
bomb() {  
    bomb | bomb &  
} bomb
```



Making the child do something

```
int main( void ) {
    if ( fork() > 0 ) {
        /* parent process */
        wait( NULL );
    } else {
        /* child process */
        execv( "/bin/bash", NULL );
    }
}
```

```
/usr/bin/zsh
 \_ ./fork_exec
      \_ [bash]
```



What does this code do?

```
int main( void ) {
    if ( fork() > 0 ) {
        /* parent process */
        sleep( 1 );
    } else {
        /* child process */
        exit( 1 );
    }
}
```



What does this code do?

```
\_ /usr/bin/zsh
```

Parent → _ ./zombie-creator

Child → _ [zombie-creator] <defunct>

Zombie Process

When a child has died but
has not been “reaped”

Child metadata stays in
process table so parent can
collect exit status

Totally normal, all children
that exit are zombies!



How to kill zombie process

I launched my program in the foreground (a daemon program), and then I killed it with `kill -9`, but I get a zombie remaining and I'm not able to kill it with `kill -9`. How to kill a zombie process?

124

share edit flag



50

edited Dec 10 '14 at 18:36



octosquidopus

1,125 ● 3 ● 20 ● 38

asked Jun 5 '13 at 16:14



MOHAMED

18.2k ● 28 ● 97 ● 169

5 Answers

active

oldest

votes



A zombie is already dead, so you cannot kill it.

168



MOHAMED

5/16/14 10:18 AM



What happens when a process dies

```
if ( fork() > 0 ) {  
    /* parent process */  
    sleep( 1 );  
} else {  
    /* child process */  
    exit( 1 );  
}
```

When a process exits,
returns int (exit code
0-255)

0 is success and anything
else indicates an error



What happens when a process dies

```
if ( fork() > 0 ) {  
    /* parent process */  
    sleep( 1 );  
} else {  
    /* child process */  
    exit( 123 );  
}
```

When a process exits,
returns int (exit code
0-255)

0 is success and anything
else indicates an error

Parents need to wait() on children

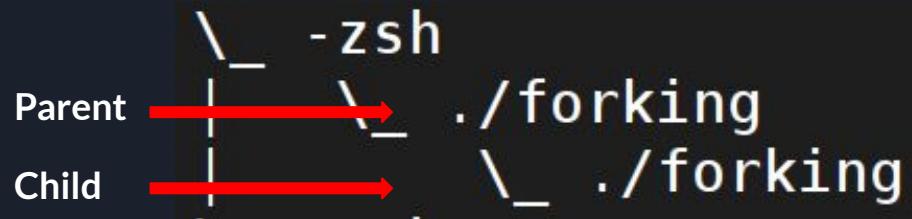
```
int main( void ) {
    if ( fork() > 0 ) {
        int status;
        sleep( 1 );
        wait( &status );
        printf( "%d\n", WEXITSTATUS( status ) );
    } else {
        exit( 123 );
    }
}
```

```
# keur @ magneto in ~/repos/temp
$ ./good-parent
123
```

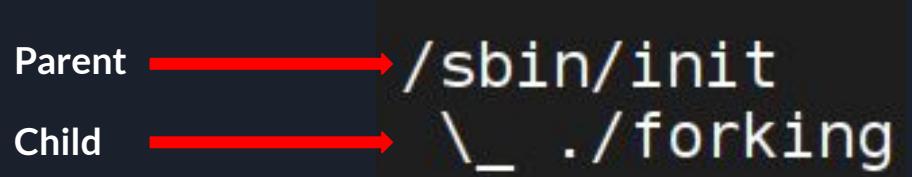
What happens if the parent exits first?

```
int main( void ) {
    if ( fork() > 0 ) {
        /* parent process */
        sleep ( 1 );
        exit( 0 );
    } else {
        /* child process */
        sleep( 100 );
        exit( 123 );
    }
}
```

Process tree after start



Process tree after 1 second



What happens if the parent exits first?

If parent exits first, **orphan** processes are re-parented by the **init** process

init reaps all orphans that are zombies





When can zombies become a problem?

Parent doesn't `wait()` on children

Parent is long running process

Zombie child processes never become orphans

Resource leakage!

Zombie leakage in the wild

```
kitty
 \_ /usr/bin/zsh
   \_ aerc
     \_ less -R
       \_ [awk] <defunct>
         \_ less -R
```

```
310bec27024579e7ada35585b3190ab875540804 – Kevin Kuehler ⚡ be4ea0d master
widgets/msgview: Reap the filter command
```

The filter command shells out and returns almost immediately. Call `Wait()` so the filter process gets reaped. Prior to this patch, `aerc` creates a zombie process for every email that is viewed.

Signed-off-by: Kevin Kuehler <keur@xcf.berkeley.edu>

```
      \_ less -R
        \_ [awk] <defunct>
          \_ less -R
            \_ [awk] <defunct>
              \_ less -R
                \_ [awk] <defunct>
                  \_ less -R
                    \_ [awk] <defunct>
```

2020 update

```
12:33 <cooperc> but what about aerc????  
12:35 <keur> no  
12:35 <keur> 🤦  
13:37 <keur> i can distinctly remember why  
13:37 <keur> 2019-11-21 11:33:10 keur i just found more zombies  
13:37 <keur> 2019-11-21 11:34:39 ddevault you shouldn't talk about your mom that way  
13:37 <keur> goodbye aerc
```



Inter-process Communication

Various ways in which processes can communicate

- Exit codes
- Signals (e.g. SIGTERM, SIGKILL, SIGINT)
- Pipes (STDIN, STDOUT, STDERR)
- Sockets (UNIX socket, IP socket)
- Message Bus (e.g. dbus on Linux)
- ... and many many more...



Process Signals

- **SIGTERM**: tell a process to exit now
- **SIGKILL**: terminate process immediately
- **SIGINT**: interrupt, when you press Ctrl+C
- **SIGHUP**: the user closes the terminal window
- **SIGWINCH**: terminal window resized
- **SIGSTOP / SIGCONT**: stop/resume

SIGKILL and SIGSTOP cannot be handled

Working with processes: ps

ps aux

```
# keur @ magneto in ~ [12:12:32]
$ ps aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START  TIME COMMAND
root      1  0.0  0.0 187592 10532 ?        Ss Oct15  0:15 /sbin/init
root      2  0.0  0.0      0     0 ?        S Oct15  0:00 [kthreadd]
root      3  0.0  0.0      0     0 ?        I< Oct15  0:00 [rcu_gp]
root      4  0.0  0.0      0     0 ?        I< Oct15  0:00 [rcu_par_gp]
root      6  0.0  0.0      0     0 ?        I< Oct15  0:00 [kworker/0:0H-kblockd]
root      8  0.0  0.0      0     0 ?        I< Oct15  0:00 [mm_percpu_wq]
root      9  0.0  0.0      0     0 ?        S Oct15  0:01 [ksoftirqd/0]
root     10  0.0  0.0      0     0 ?        I Oct15  0:14 [rcu_preempt]
root     11  0.0  0.0      0     0 ?        S Oct15  0:03 [rcuc/0]
root     12  0.0  0.0      0     0 ?        S Oct15  0:00 [rcub/0]
root     13  0.0  0.0      0     0 ?        S Oct15  0:00 [migration/0]
root     14  0.0  0.0      0     0 ?        S Oct15  0:00 [idle_inject/0]
root     16  0.0  0.0      0     0 ?        S Oct15  0:00 [cpuhp/0]
root     17  0.0  0.0      0     0 ?        S Oct15  0:00 [cpuhp/1]
root     18  0.0  0.0      0     0 ?        S Oct15  0:00 [idle_inject/1]
root     19  0.0  0.0      0     0 ?        S Oct15  0:00 [migration/1]
root     20  0.0  0.0      0     0 ?        S Oct15  0:02 [rcuc/1]
root     21  0.0  0.0      0     0 ?        S Oct15  0:01 [ksoftirqd/1]
root     23  0.0  0.0      0     0 ?        I< Oct15  0:00 [kworker/1:0H-kblockd]
root     24  0.0  0.0      0     0 ?        S Oct15  0:00 [cpuhp/2]
root     25  0.0  0.0      0     0 ?        S Oct15  0:00 [idle_inject/2]
root     26  0.0  0.0      0     0 ?        S Oct15  0:00 [migration/2]
root     27  0.0  0.0      0     0 ?        S Oct15  0:03 [rcuc/2]
root     28  0.0  0.0      0     0 ?        S Oct15  0:01 [ksoftirqd/2]
root     30  0.0  0.0      0     0 ?        I< Oct15  0:00 [kworker/2:0H-kblockd]
root     31  0.0  0.0      0     0 ?        S Oct15  0:00 [cpuhp/3]
root     32  0.0  0.0      0     0 ?        S Oct15  0:00 [idle_inject/3]
root     33  0.0  0.0      0     0 ?        S Oct15  0:00 [migration/3]
```

Working with processes: ps

ps fwaux

```
# keur @ magneto in ~ [12:17:28]
$ ps fwaux | tail -n 10
keur      70404  0.1  0.0  13448  9356 pts/16    Ss+  10:33   0:10  \_ /usr/bin/zsh
keur      80593  0.1  0.2 1014872 33864 ?          Sl   11:45   0:02  alacritty
keur      80597  0.0  0.0  12752  8604 pts/20    Ss   11:45   0:00  \_ /usr/bin/zsh
keur      80747  0.0  0.0  10148  4708 pts/20    S+   11:46   0:00      \_ man pstree
keur      80757  0.0  0.0   6504  2684 pts/20    S+   11:46   0:00      \_ less
root     81103  0.0  0.0 310816  3636 ?          Ssl  12:09   0:00 /usr/bin/pcscd --foreground --auto-exit
keur      81634  0.8  0.6 276708 98572 ?          Sl   12:12   0:02  kitty
keur      81638  0.7  0.0  12728  8668 pts/1     Ss   12:12   0:02  \_ /usr/bin/zsh
keur      82142  0.0  0.0  13184  3564 pts/1     R+   12:17   0:00      \_ ps fwaux
keur      82143  0.0  0.0   5332   760 pts/1     S+   12:17   0:00      \_ tail -n 10
```

Working with processes: ps

```
ps -eo pid,user,cmd --sort user
```

```
# keur @ magneto in ~ [12:20:10]
$ ps -eo pid,user,cmd --sort user | head -n 18
  PID USER      CMD
  537 colord    /usr/lib/colord
  505 dbus      /usr/bin/dbus-daemon --system --address=systemd: --nofork
  803 keur      /usr/lib/systemd/systemd --user
  804 keur      (sd-pam)
  809 keur      /usr/bin/ssh-agent -D -a /run/user/1000/ssh-agent.socket
  817 keur      /usr/lib/gdm-wayland-session --register-session sway
  819 keur      /usr/bin/dbus-daemon --session --address=systemd: --nofork
  821 keur      sway
  829 keur      swaybg -o eDP-1 -i /home/keur/Pictures/purple.png -m fill
l -o DP-2 -i /home/keur/Pictures/miku1.jpg -m fill
  832 keur      swaybar -b bar-0
  834 keur      mako
  837 keur      i3status
  891 keur      alacritty
  895 keur      /usr/bin/zsh
  903 keur      /usr/bin/gpg-agent --supervised
  948 keur      mosh-client -# keur@23.239.23.104 | 23.239.23.104 60015
  954 keur      scdaemon --multi-server
```

Working with processes: ps

pstree



Working with processes: kill, pkill

Send TERM to process

- \$ kill 100

Reload SSH process

- \$ pkill -HUP sshd

Send SIGKILL to chrome

- \$ pkill -9 google-chrome-stable

Working with processes: pgrep

```
# keur @ magneto in ~ [13:08:13]
$ pgrep brave
69169
69172
69174
69195
69201
69370
69371
69376
69412
69434
69439
69446
69609
70623
70703
72760
74619
```

Good for command
chaining with xargs



Working with processes: pgrep

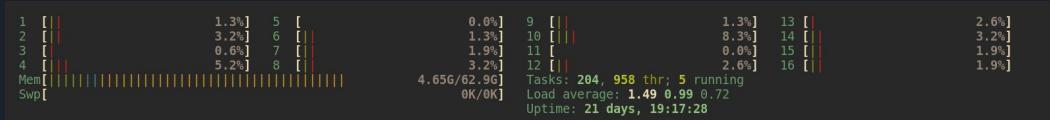
Mosh servers stay around waiting for connections

I know these connections are dead

- `$ pgrep mosh-server | grep -v $PPID | xargs kill`

Kill mosh servers except the current connection

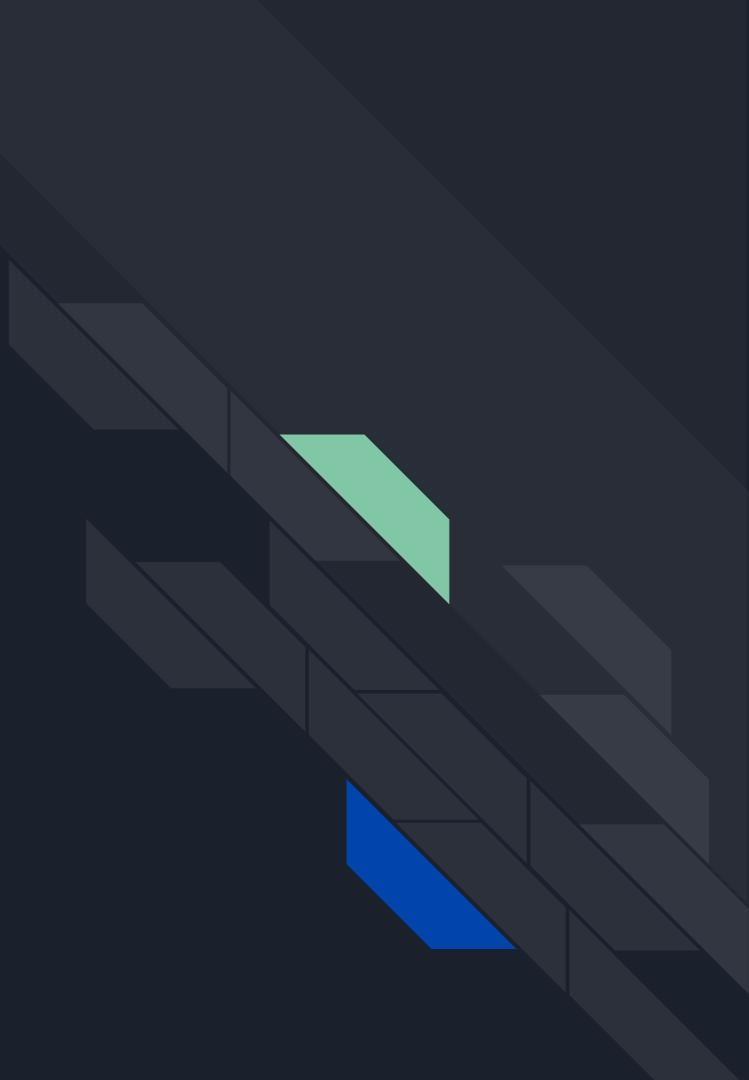
Working with processes: htop



PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
7077	root	20	0	427M	18676	13196	R	6.5	0.0	0:00.19	/usr/sbin/runc init
26992	root	20	0	473M	298M	72972	S	3.9	0.5	2h11:26	kube-apiserver --advertise-address=23.169.96.55 --allow-privileged=true --authoriz
4521	root	20	0	539M	335M	73120	S	3.9	0.5	1h16:04	kube-apiserver --advertise-address=23.169.96.55 --allow-privileged=true --authoriz
6634	root	20	0	33384	6020	3920	R	3.2	0.0	0:00.66	htop
26910	root	20	0	219M	101M	62196	S	3.2	0.2	58:53.16	kube-controller-manager --allocate-node-cidrs=true --authentication-kubeconfig=/etc
3361	root	20	0	3192M	101M	68348	S	2.6	0.2	5:15.95	/usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --k
25581	root	20	0	3263M	99828	67872	S	1.9	0.2	1h12:39	/usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --k
24392	root	20	0	473M	315M	72876	S	1.9	0.5	2h08:53	kube-apiserver --advertise-address=23.169.96.55 --allow-privileged=true --authoriz
4615	root	20	0	10.16	72448	22264	S	1.9	0.1	30:07.89	etcd --advertise-client-urls=https://23.169.96.55:2379 --cert-file=/etc/kubernetes
22330	root	20	0	3119M	99228	67372	S	1.9	0.2	1h12:23	/usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --k
4994	root	20	0	219M	98M	62200	S	1.9	0.2	32:00.01	kube-controller-manager --allocate-node-cidrs=true --authentication-kubeconfig=/etc
26670	root	20	0	10.16	69752	21448	S	1.9	0.1	53:21.53	etcd --advertise-client-urls=https://23.169.96.55:2379 --cert-file=/etc/kubernetes
24651	root	20	0	219M	100M	62068	S	1.9	0.2	59:11.45	kube-controller-manager --allocate-node-cidrs=true --authentication-kubeconfig=/etc
24321	root	20	0	10.16	69812	21444	S	1.3	0.1	54:08.12	etcd --advertise-client-urls=https://23.169.96.55:2379 --cert-file=/etc/kubernetes
4838	root	20	0	539M	335M	73120	S	1.3	0.5	13:36.66	kube-apiserver --advertise-address=23.169.96.55 --allow-privileged=true --authoriz
13565	root	20	0	2615M	1129M	18616	S	1.3	1.8	14h12:30	/usr/bin/qemu-system-x86_64 -name deus -nographic -m 2048 -drive file=/dev/ubuntu-
5092	root	20	0	539M	335M	73120	S	1.3	0.5	1:54.54	kube-apiserver --advertise-address=23.169.96.55 --allow-privileged=true --authoriz
29643	root	20	0	219M	101M	62196	S	1.3	0.2	1:21.39	kube-controller-manager --allocate-node-cidrs=true --authentication-kubeconfig=/etc
3498	root	20	0	3192M	101M	68348	S	1.3	0.2	1:36.39	/usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --k
5013	root	20	0	10.16	72448	22264	S	1.3	0.1	0:51.50	etcd --advertise-client-urls=https://23.169.96.55:2379 --cert-file=/etc/kubernetes
11053	root	20	0	146M	38084	28708	S	1.3	0.1	0:08.75	/coredns -conf /etc/coredns/Corefile
13589	root	20	0	2615M	1129M	18616	S	0.6	1.8	13h24:52	/usr/bin/qemu-system-x86_64 -name deus -nographic -m 2048 -drive file=/dev/ubuntu-
31722	root	20	0	3614M	52144	31980	S	0.6	0.1	9:52.20	/usr/local/bin/crio
27180	root	20	0	473M	298M	72972	S	0.6	0.5	25:22.14	kube-apiserver --advertise-address=23.169.96.55 --allow-privileged=true --authoriz
24923	root	20	0	219M	100M	62068	S	0.6	0.2	2:32.31	kube-controller-manager --allocate-node-cidrs=true --authentication-kubeconfig=/etc
24743	root	20	0	473M	315M	72876	S	0.6	0.5	3:13.01	kube-apiserver --advertise-address=23.169.96.55 --allow-privileged=true --authoriz
27200	root	20	0	473M	298M	72972	S	0.6	0.5	3:04.54	kube-apiserver --advertise-address=23.169.96.55 --allow-privileged=true --authoriz

- Aesthetically pleasing
- Interactive (curses)
- Does what all the other tools can do

Part 2: Init and Services





Two types of processes

Not strict definitions...

Foreground process: chrome, vim, htop.

Started and stopped by the user.

Daemons: background processes like
sshd, nginx, postfix, etc.



What is a service?

A service is a **daemon** managed by your init system

Usually a collection of processes working together



How do we control processes

Many ways...

- We can send signals to the process
- We can use CLI tools (apachectl, prosodyctl)
- We can have the init system do the above



What is an init system?

The daemon that manages all other daemons

First process started by the kernel (PID 1)

Starts enough processes to make the system useful

Stops processes on shutdown

Reaps orphans that are zombies



Bare bones init system

Have the user manage the system

Okay, let's just have `init` spawn a shell

Bare bones init system

```
pid_t result;  
if ( result = fork() > 0 ) {  
    /* parent process */  
    while ( wait ( NULL ) != result ) { }  
} else {  
    /* child process */  
    execv( "/bin/bash", NULL );  
}
```

Reap all processes

Start other processes



Traditional init systems

Other responsibilities

- Start and stop daemons at certain times (runlevels) during the boot chain
- Launch services
- Send signals to services



How does a traditional init system do it?

There is a script for each service

```
/etc/init.d/$service $action
```

Each script acts on one process



Why are actions better than manually using tools?

Instead of figuring out how to signal every service on my system, I remember 5 actions that apply to *all* services.

If I want details, I can look at the scripts.



What is a “modern” init system

Traditional	Modern
Shell scripts that can do <i>anything</i>	Declarative configuration files
Stores state in PID files	Stores state in init system
Actions can be customized per script	Consistent actions provided by CLI tool
Difficult to control ordering	Ordering by dependencies or events

A traditional init script - NGINX

```
1 start() {
2     [ -x $nginx ] || exit 5
3     [ -f $NGINX_CONF_FILE ] || exit 6
4     make_dirs
5     echo -n $"Starting $prog: "
6     daemon $nginx -c $NGINX_CONF_FILE
7     retval=$?
8     echo
9     [ $retval -eq 0 ] && touch $lockfile
10    return $retval
11 }
12 stop() {
13     echo -n $"Stopping $prog: "
14     killproc $prog -QUIT
15     retval=$?
16     echo
17     [ $retval -eq 0 ] && rm -f $lockfile
18     return $retval
19 }
20
21 restart() {
22     configtest || return $?
23     stop
24     sleep 1
25     start
26 }
27
28 reload() {
29     configtest || return $?
30     echo -n $"Reloading $prog: "
31     killproc $prog -HUP
32     retval=$?
```

- Only a subsection of the script
- Tedious to write
- Lots of copy-pasted code
- Keeps state using lock files and PID files - racey!



Modern init system - systemd

Developed at Red Hat

Based off macOS launchd

Heavily parallelized



Dependency and event based

systemd service - NGINX

```
[Unit]
Description=A high performance web server and a reverse proxy server
Documentation=man:nginx(8)
After=network.target

[Service]
Type=forking
PIDFile=/run/nginx.pid
ExecStartPre=/usr/sbin/nginx -t -q -g 'daemon on; master_process on;'
ExecStart=/usr/sbin/nginx -g 'daemon on; master_process on;'
ExecReload=/usr/sbin/nginx -g 'daemon on; master_process on;' -s reload
ExecStop=-/sbin/start-stop-daemon --quiet --stop --retry QUIT/5 --pidfile /run/nginx.pid
TimeoutStopSec=5
KillMode=mixed

[Install]
WantedBy=multi-user.target
```

Anatomy of a service file

/etc/systemd/system/my-cool-server.service

[Unit]

Description=My cool server

Dependency on network

Requires=network-online.target

After=network-online.target

Run our executable

[Service]

ExecStart=/usr/local/bin/my-cool-server

When service can start automatically

[Install]

WantedBy=multi-user.target

systemctl (no arguments)

getty@tty1.service	loaded active running	Getty on tty1
haveged.service	loaded active running	Entropy daemon using the HAVEGE algorithm
irqbalance.service	loaded active running	irqbalance daemon
keyboard-setup.service	loaded active exited	Set the console keyboard layout
kmmod-static-nodes.service	loaded active exited	Create list of required static device nodes
loadcpufreq.service	loaded active exited	LSB: Load kernel modules needed to enable c
lvm2-lvmetad.service	loaded active running	LVM2 metadata daemon
lvm2-monitor.service	loaded active exited	Monitoring of LVM2 mirrors, snapshots etc.
munin-node.service	loaded active running	Munin Node
netfilter-persistent.service	loaded active exited	netfilter persistent configuration
networking.service	loaded active exited	Raise network interfaces
nginx.service	loaded active running	A high performance web server and a reverse
node_exporter.service	loaded active running	Prometheus node_exporter
ntp.service	loaded active running	LSB: Start NTP daemon
postfix.service	loaded active exited	Postfix Mail Transport Agent
postfix@-.service	loaded active running	Postfix Mail Transport Agent (instance -)
quota.service	loaded active exited	Initial Check File System Quotas
rpc-gssd.service	loaded active running	RPC security service for NFS client and ser
● rpc-svcgssd.service	loaded failed failed	RPC security service for NFS server
rpcbind.service	loaded active running	RPC bind portmap service
rsyslog.service	loaded active running	System Logging Service
serial-getty@ttyS0.service	loaded active running	Serial Getty on ttyS0
ssh.service	loaded active running	OpenBSD Secure Shell server
sysstat.service	loaded active exited	LSB: Start/stop sysstat's sadc
systemd-journal-flush.service	loaded active exited	Flush Journal to Persistent Storage
systemd-journald.service	loaded active running	Journal Service
systemd-logind.service	loaded active running	Login Service
systemd-modules-load.service	loaded active exited	Load Kernel Modules
systemd-random-seed.service	loaded active exited	Load/Save Random Seed
systemd-remount-fs.service	loaded active exited	Remount Root and Kernel File Systems
systemd-sysctl.service	loaded active exited	Apply Kernel Variables
systemd-tmpfiles-setup-dev.service	loaded active exited	Create Static Device Nodes in /dev
systemd-tmpfiles-setup.service	loaded active exited	Create Volatile Files and Directories
systemd-udev-trigger.service	loaded active exited	udev Coldplug all Devices
systemd-udevd.service	loaded active running	udev Kernel Device Manager
systemd-update-utmp.service	loaded active exited	Update UTMP about System Boot/Shutdown
systemd-user-sessions.service	loaded active exited	Permit User Sessions

`systemctl status $service`

```
# keur @ supernova in ~ [15:59:50]
$ systemctl status nginx
● nginx.service - A high performance web server and a reverse proxy server
  Loaded: loaded (/lib/systemd/system/nginx.service; enabled; vendor preset: enabled)
  Active: active (running) since Fri 2019-10-11 14:26:13 PDT; 5 days ago
    Docs: man:nginx(8)
 Process: 809 ExecStart=/usr/sbin/nginx -g daemon on; master_process on; (code=exited, status=0/SUCCESS)
 Process: 725 ExecStartPre=/usr/sbin/nginx -t -q -g daemon on; master_process on; (code=exited, status=0/SUCCESS)
 Main PID: 837 (nginx)
   Tasks: 9 (limit: 4915)
  Memory: 20.9M
     CPU: 1min 53.810s
```

`systemctl start $service`

`systemctl enable $service`



journalctl

Log aggregator built into systemd

View recent logs for a unit

- `$ journalctl -exu $service`

Tail logs (`tail -f /var/log/nginx.log`)

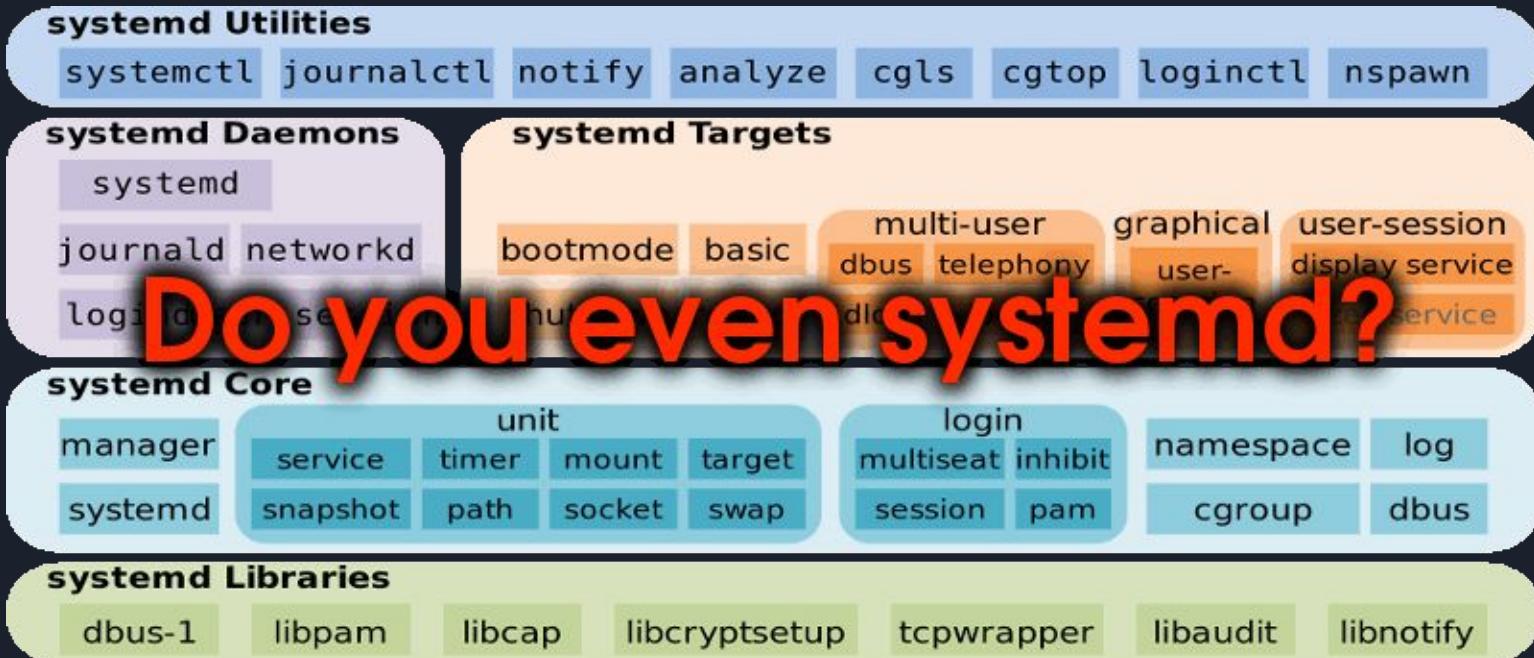
- `$ journalctl -fu $service`

journalctl - Declarative logs

```
$ journalctl --since today --until now
```

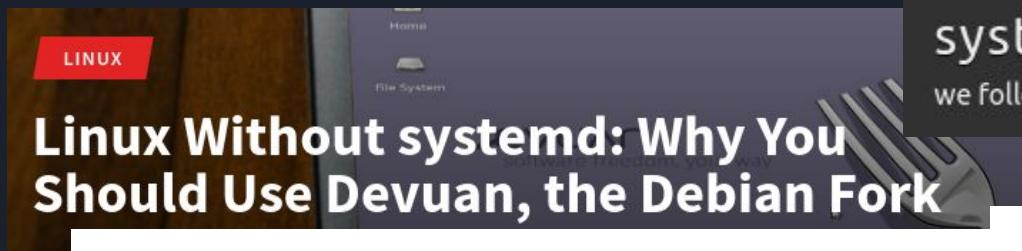
```
# keur @ magneto in /lib/systemd/system [16:11:12]
$ journalctl --since today --until "3 hours ago" | head -n 15
-- Logs begin at Wed 2018-09-26 21:48:11 PDT, end at Wed 2019-10-16 16:09:05 PDT. --
Oct 16 00:28:42 magneto kernel: Filesystems sync: 0.013 seconds
Oct 16 00:28:42 magneto kernel: Freezing user space processes ... (elapsed 0.007 seconds) done.
Oct 16 00:28:42 magneto kernel: OOM killer disabled.
Oct 16 00:28:42 magneto kernel: Freezing remaining freezable tasks ... (elapsed 0.001 seconds) done.
Oct 16 00:28:42 magneto kernel: printk: Suspending console(s) (use no_console_suspend to debug)
Oct 16 00:28:42 magneto kernel: wl0: deauthenticating from c0:8a:de:30:94:dc by local choice (Reason: 3=DEAUTH_LEAVING)
Oct 16 00:28:42 magneto kernel: sd 0:0:0:0: [sda] Synchronizing SCSI cache
Oct 16 00:28:42 magneto kernel: sd 0:0:0:0: [sda] Stopping disk
Oct 16 00:28:42 magneto kernel: e1000e: EEE TX LPI TIMER: 00000011
Oct 16 00:28:42 magneto kernel: ACPI: EC: interrupt blocked
Oct 16 00:28:42 magneto kernel: ACPI: Preparing to enter system sleep state S3
Oct 16 00:28:42 magneto kernel: ACPI: EC: event blocked
Oct 16 00:28:42 magneto kernel: ACPI: EC: EC stopped
Oct 16 00:28:42 magneto kernel: PM: Saving platform NVS memory
```

systemd is really complicated



We barely scratched the surface...

systemd is really controversial



systemd-free linux community

we follow the development of linux away from systemd

THIS IS Linux

THIS IS Linux without systemd

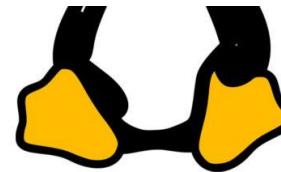
systemd is the best example of Suck.

There is a menace which is spreading like a disease throughout the Linux world, it is called systemd.

Systemd Sucks
@systemdsucks

I only RT if in eng. Helping sysadmins vent since 2016. FOSS diversity initiative community building effort going on: /join #systemdsucks on freenode!

pid1 Joined July 2016





systemd goals

- Performance
- Simple UX
- Providing “building blocks” for an operating system

Ended up re-implementing many existing tools along the way...



My take...

systemd is easy to use for simple things

Fast init systems are actually pretty complicated

At this point, you don't really have a choice
anyway...

its not dns