



# Config Management

Ja (Thanakul) Wattanawong (jaw)

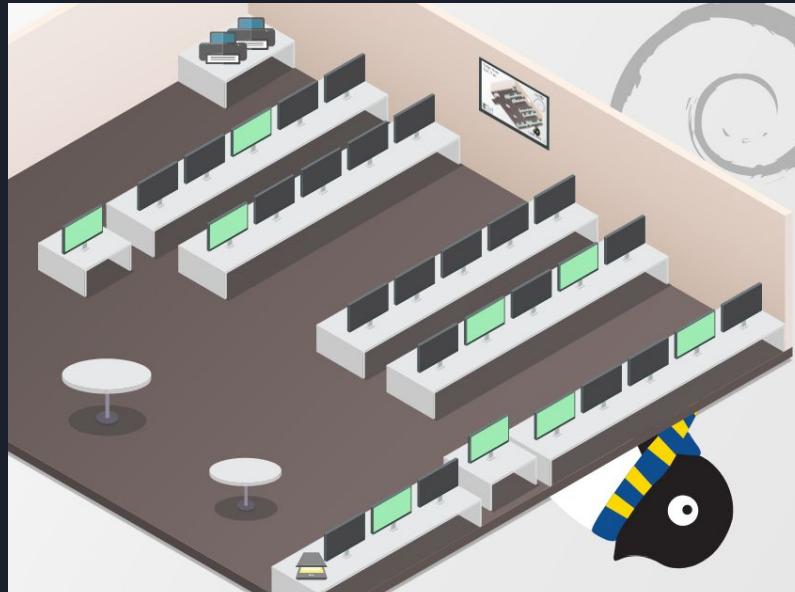


# Who am I?

- OCF Site Manager
- EECS
- Firefighter
- I set up your student VMs
- Catch me in the ~~server room~~ OCF social discord

# What problems does config management solve (1/3)?

- Suppose you have a bunch of computers



- You suddenly decide that everybody computer in the lab needs Minecraft installed
- Without config management: SSH into all the desktops and install it
- Even with a script this sucks
- Problem: How do you deploy updates to a fleet of existing computers?

# What problems does config management solve (2/3)?

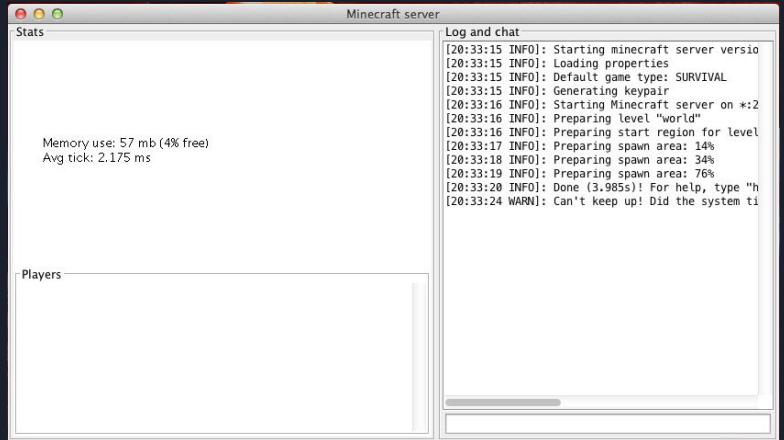
- Suppose this computer lab buys a new computer



- Remembering to install Minecraft and configure it correctly is pretty difficult
  - Eg: Have it use more than 1Gb RAM by default
- Problem: How do you provision new machines?

# What problem does config management solve (3/3)?

- Suppose you are running a Minecraft Server.



- You realize some change you made long ago broke some minor thing
- How do you figure out what settings you had changed around that time?
- Problem: How do you communicate what changes in software configuration to future you (and others)?

# Config Management

✓ Add minecraft-launcher to packages.pp

by master

 fydai committed on Feb 12 Verified

1 parent be4bc6a commit 466fc20ef26470425792c5d9df4bc03ba21893c8

Showing 1 changed file with 1 addition and 1 deletion.

Unified Split

modules/ocf\_desktop/manifests/packages.pp

```
diff --git a/modules/ocf_desktop/manifests/packages.pp b/modules/ocf_desktop/manifests/packages.pp
--- a/modules/ocf_desktop/manifests/packages.pp
+++ b/modules/ocf_desktop/manifests/packages.pp
@@ -29,7 +29,7 @@ 
 29 29      # FUSE
 30 30      ['fuse', 'exfat-fuse'];
 31 31      # games
- 32 -      ['armagetronad', 'gl-117', 'gnome-games', 'wesnoth', 'wesnoth-music'];
+ 32 +      ['armagetronad', 'gl-117', 'gnome-games', 'minecraft-launcher', 'wesnoth', 'wesnoth-music'];
 33 33      # graphics/plotting
 34 34      ['r-cran-rgl', 'jupyter-qtconsole', 'rstudio'];
 35 35      # input method editors
```



# Config management

- Solves problem 1 (updating computers) by having an unified update mechanism
- Solves problem 2 (bootstrapping new computers) by having all the changes necessary in some centralized repository
- Solves problem 3 (communicating new changes) by allowing you to use standard development practices (mainly git) to record your changes, and communicate with others

Configuration Management - Software that makes it as easy as possible to bootstrap new machines, configure running software, and allows configuration to be stored as code  
“configuration as code” philosophy



# Configuration Management Philosophies

- Imperative:
  - Treats configuration as a “set of tasks”, order to be specified by you.
  - Say “How you want to do it”
  - “Install minecraft”, then “add a line to the config file”, then “run minecraft”
  - Examples: Chef, Ansible
  - Updates handled differently than Bootstrapping
  - What if config file is already edited?
- Declarative
  - Specify the final state, the system works to get itself into the state
  - Say “What you want, software figures out how to do it”
  - “Ensure minecraft is installed, the config file has line <X> in it, and ensure that minecraft is running”
  - Examples: Puppet
  - Updates are handled the same as Bootstrapping

Of course, this is not an either-or, any software will have aspects of both philosophies



# Puppet

- Popular Configuration Management software
- Used for configuring individual machines
- Declarative philosophy, with some Imperative components when necessary
- Originally built on Ruby, now its own configuration language
- Used at places like
  - OCF (<https://github.com/ocf/puppet>)
  - CS 162 (<https://github.com/Berkeley-CS162/vagrant/tree/master/modules/cs162>)
  - Wikimedia (<https://github.com/wikimedia/puppet>)
  - Github
  - Lyft
- “Pull model” - Configured machines ask for an update
  - So Puppet is usually scheduled to be run every now and then (OCF has 30 minutes)

# What happens when Puppet is run?

- Client asks server for an update
  - “I want to be configured as a Minecraft Server”
- Server asks client for a list of Facts
  - “Ok, send me your hostname, and RAM”
- Client responds with the facts
  - “My hostname is zombies.ocf.berkeley.edu and I have 4GB RAM”
- Server responds with configuration
  - “Ensure the Minecraft server is running, with hostname zombies.ocf.berkeley.edu, 4GB RAM, with this configuration file”
- Client makes the necessary changes to ensure its current configuration matches the configuration given by the server
  - “The minecraft server is currently running, but the configuration file has been updated, I will fetch the updated version”

## Lifecycle of a Puppet Run





# Puppet Code

- Most of the code is here
  - Files - contains static files
  - Templates - contain templates (Ruby style)
  - Manifests - the heart of the configuration, specifies the desired states
- Other sections that are occasionally used
  - Facts - Ways to extract data needed for configuration
  - Functions - if you need extra something fancy data structure manipulating
- Dependencies need to be explicitly described
  - Puppet is allowed to run code in any order that it sees fit
  - If you have code installing Minecraft, and running Minecraft, you need to tell puppet to install Minecraft before running it



# Example Puppet Code - Adding a user and a home directory

```
user { 'ocftv':
  comment => 'TV NUC',
  home    => '/opt/tv',
  groups   => ['sys', 'audio'],
  shell    => '/bin/bash';
}

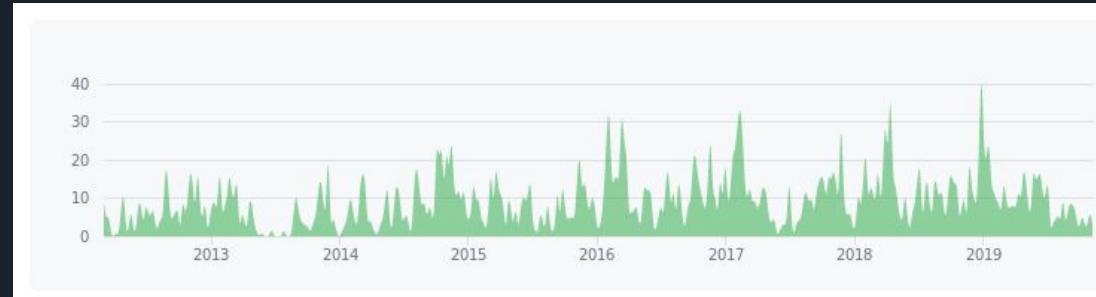
file {
  # Create home directory for ocftv user
  '/opt/tv':
    ensure  => directory,
    owner   => ocftv,
    group   => ocftv,
    require => User['ocftv'];
}
```

# Example Puppet Code - Running a web server

```
1 package { 'nginx'; }
2 service { 'nginx':
3   require  => Package['nginx'],
4   subscribe => Class['ocf::ssl::default'],
5 }
6
7 file {
8   '/etc/nginx/conf.d/local.conf':
9   content => template('ocf_apphost/local.conf.erb'),
10  require => Package['nginx'],
11  notify  => Service['nginx'];
12
13 # raise the hash bucket size for server names since we use really long server
14 # names (like something.apphost.ocf.berkeley.edu)
15 #
16 # http://nginx.org/en/docs/http/server_names.html
17 server_names_hash_bucket_size 128;
18
19 ssl_dhparam /etc/ssl/dhparam.pem;
20 ssl_protocols <%= @ssl_protocols %>;
21 ssl_ciphers '<%= @ssl_ciphersuite %>';
22
23 # combined log format, with virtual host added (rt#4459)
24 log_format vhost '$host $remote_addr - $remote_user [$time_local] '
25           '$request" $status $body_bytes_sent '
26           '"$http_referer" "$http_user_agent"';
27
28 # increase client max request body size (default is 1MiB)
29 client_max_body_size 20M;
```

# Puppet at the OCF (1/2)

<https://github.com/ocf/puppet>



- Originally started in 2012, from the “Configuration is edited directly on the server, and desktops manually” model
  - We only had 10 desktops so this was kinda okay
- 7 years later, all of the OCF’s machines runs off the puppet repository
  - Desktops
  - Thing behind the TV
  - Hypervisors (things running the VMs)
  - VMs (Running all the Networked Services you learned about)
    - Including the puppet server itself



# Puppet at the OCF (2/2)

- All the code is split into modules
  - Ocf\_tv
  - Ocf\_desktop
  - Ocf\_www
  - Ocf\_printhost
- Common OCF modules for shared configuration
  - Ocf::ssl for (I need a web certificate)
  - Ocf::auth for LDAP and Kerberos and sudoers configuration

# Bonus Slide: Terraform

- Has integrated APIs to provision machines declaratively on cloud platforms
- This is part of the code used to generate your decal VMs. Notice that the provisioning script is mostly imperative.
- The alternative (which we seriously considered) was clicking “New droplet” 80 times.

```
54   provisioner "remote-exec" {
55     connection {
56       type    = "ssh"
57       user    = "root"
58       private_key = "${file("${var.pvt_key_file}")}"
59       host    = self.ipv4_address
60     }
61     inline = [
62       # Set the hostname to be FQDN
63       "sudo hostname ${each.value.username}.decal.xcf.sh",
64
65       # Add the user and give them root access
66       "sudo useradd ${each.value.username} -s /bin/bash -m",
67       "sudo echo ${each.value.username}:${each.value.password} | sudo chpasswd",
68       "sudo usermod -aG sudo ${each.value.username}",
69       "sudo service sshd restart",
70       "sudo passwd -e ${each.value.username}",
71
72       # Populate the motd with data
73       "sudo sed -i 's/$HOSTNAME/${each.value.username}/g' /etc/motd",
74       "sudo sed -i 's/$IP/${self.ipv4_address}/g' /etc/motd",
75
76       # Reboot to allow MOTD to change. This is a hack.
77       "sudo shutdown -r +60"
78     ]
79   }
```



Thank you for your time!

The lab is up!

If you have any issues, or the lab is unclear, please don't hesitate to ask me questions! Also ask in #decal-general