

Containers, Container Orchestration, Architecture

Benjamin Zhang <bzh>
Based on Brian Sang's Fall 2019 slides

Where we left off

- At scale, failure is common
- Solved problem of configuring large number of machines


This week

- How to make our applications run on large systems?
 - Seamless handover if a server keels over
 - How does Google/FB/etc. have nearly perfect uptime?

Containers


What are VMs exactly?

- Virtual Machines
- Emulates a physical machine using software
- Typical setup: hypervisor runs multiple VMs
- Provides isolation
- Some bloat: Need different Guest OS and emulation of Virtual Hardware for each application
- Takes some time to boot up




What is a process, exactly?

- Take CS162
 - Best class
 - Don't argue with me
- An executing program
- Independent virtual address space, program, CPU state, program state, file descriptor #s...
- ...and more. Everything you need to run a program (shell, web browser, pubg, etc.)
 - Includes other binaries and libraries which may be dynamically linked




Containers - between VMs and processes?




Containers are isolated,
but share OS and, where
appropriate, bins/libraries

...result is significantly faster deployment,
much less overhead, easier migration,
faster restart




Containers vs. Processes


textbook process



real process




container



Containers — Big Picture!

- Often compared with VMs, but are more like processes with bundled environments
- Provide similar isolation
 - Much less than VMs, however! For this reason, we often run containers inside VMs still (but can run >1 container per VM)
- Much faster to boot up than VMs
- Easy to package applications into containers
- Distribute built containers (*images*) to quickly deploy
- Common container runtimes: Docker, rkt, LXC




Docker Concepts - Images and Containers

- **Image:** an executable package that includes everything needed to run an application
 - the code, a runtime, libraries, environment variables, and configuration files.
- **Container:** A runnable/running instance of an image
- Similar to Class vs. Object concept in OOP
- Can pull/push built images from a registry, e.g. [DockerHub](#)
 - Or a private registry

Docker Concepts - Building an Image, Dockerfiles


- Images need to be built
- Typically use a **Dockerfile** to specify how to build an image
- Images are built in layers
 - Like an onion :shrek:
 - Allows for images based off the same layer to be built faster
- Keep each layer minimal



Dockerfile Example

```
1 FROM ubuntu:xenial ← Use Ubuntu as base image
2
3 RUN apt-get update && apt-get install -y \
4     git \
5     python3 \
6     --no-install-recommends && rm -rf /var/lib/apt/lists/* ← Update the system and install
                                                               requisite packages
                                                               (delete extra files to keep things
                                                               minimal!)
7 RUN pip install Flask
8
9 ADD hello.py /app/hello.py ← Add local files to run the
                               server to our container
10 WORKDIR /app
11
12 EXPOSE 5000
13 ENV FLASK_APP=hello.py ← Run the webserver!
14 CMD ["flask", "run", "host", "0.0.0.0"]
```

The Docker daemon



Container Orchestration

Container Orchestrators


- So we have the capability of putting all of our apps into nice containers, how do we actually use them?
- Container Orchestrators help us make sure we have the right amount of containers we want, that they are on the proper machines, restart them when they die to handle failure, etc.
- Use **Distributed Systems Magic™** to recover from failures; multiple masters, distributed key/value stores, and more
- Examples: [Kubernetes](#), [Mesos+Marathon](#), [Docker Swarm](#)



kubernetes



MESOS MARATHON




A taste of Distributed Systems Magic™

Typical Structure - Master/Worker Architecture

- Multiple masters in case one fails
 - Maintain quorum using etcd, Zookeeper (distributed key/value stores)
 - Use Paxos, Raft to maintain consensus (badass two phase commit)
- Some sort of scheduler
 - Place containers on machines
- Some sort of cluster manager
 - Scale up/down machines and have them join the cluster

Mesos example




Distributed Systems is hard

- Need to make sure all nodes are actually doing what a master/coordinator says
 - Problem: a node could fail and stop responding at any time!
 - Could fail right after it says “Yes, I’m about to do this”
- All of these algorithms are just for getting nodes to AGREE on something
 - 2PC - too slow
 - Paxos - notoriously hard to understand
 - Raft - supposedly easier to understand. I don’t really understand it (haven’t read paper yet lol)
- Don’t roll your own crypto? Don’t write your own consensus algorithm
 - Just use etcd (raft) or zookeeper (custom algorithm, zab)


Architecture and Design

Don't read ahead! This part is meant to
be interactive


Basic File storage service



Divide into separate services




Improvement 1: More handlers




- If one of the handlers fails, now have more!
- Helps distribute load amongst handlers - scale out
- New problem: Database disk I/O is maxing out, what do we do now?

Improvement 2: Caching!




- Use Redis or Memcached (keep data in memory, fast)
- Keep caches as close to the client as possible - locality
- Another problem to fix: what if database goes down?

Improvement 3: Replication



- Same idea as having a backup
- Can use replica for reads if low latency
- Problems, again: Your app is now on front page of Reddit and the disk that your database uses is full

Improvement 4: Partition



- Split database into multiple machines
- Need to route requests to correct DB
- More complexity so that's kinda bad
- Now we can scale out application servers layer and data layer

Recap

- With scale, new problems
- Tradeoffs tradeoffs tradeoffs
- Scale horizontally when possible
 - But, this introduces complexity
- UNIX philosophy: small sharp tools
 - Small, sharp services scale better
 - Introduces complexity when you need to combine them