

The Multiplicative Weights Algorithm

David Woodruff

The Experts Problem

- n “experts” predict an outcome on each day
- Expert = someone with an opinion, not necessarily someone who knows anything
- For example, the experts could try to predict the stock market

Expt 1	Expt 2	Expt 3	neighbor's dog	truth
down	up	up	up	up
down	up	up	down	down
...

The Experts Problem

- n “experts” predict an outcome on each of T days, $t = 1, \dots, T$
- On day t , the i -th expert predicts outcome out_i^t
- On day t , you see $\text{out}_1^t, \dots, \text{out}_n^t$ and make your prediction guess^t
- Then you see the actual outcome out^t on day t
- You are correct if $\text{guess}^t = \text{out}^t$ and wrong otherwise

The Experts Problem

- **Goal:** if the best expert is wrong on M days, you want to be wrong on at most M days, plus a little bit
- Don't make assumptions on the input
- Don't assume future looks like the past
- You want to do as well as the best single expert in hindsight

How should you choose your guess on each day?

Simpler Question

- Suppose at least one expert is perfect, i.e., never makes a mistake
 - Don't know which one
- Suppose each expert predicts one of two values: 0 or 1
 - Stock market will go up or down
- Can we find a strategy that makes no more than $\lfloor \lg_2 n \rfloor$ mistakes?
- **Majority-and-halving:** On each day, take the majority vote of all experts
 - Each time you're wrong, you can remove at least half the experts
 - After $\lfloor \lg_2 n \rfloor$ mistakes you're left with the perfect expert
- Same guarantee if experts predict more than 2 values
 - You choose most frequent prediction. If wrong, at least half the experts are wrong

Can You Do Better?

- **Claim:** in the worst case, any deterministic strategy makes $\lg_2 n$ mistakes
- **Proof: adversary method**
- Day 1: make the first $n/2$ experts say 0, and the second $n/2$ experts say 1
 - If predictor outputs 0, then say the best expert outputs 1
 - If predictor outputs 1, then say the best expert outputs 0
 - Perfect expert is either in $[1, n/2]$ or in $[n/2+1, n]$
- Day 2: in each interval $[1, n/2]$ and $[n/2+1, n]$, make first half of the experts say 0 and second half of the experts say 1
 - If predictor outputs 0, then say the best expert outputs 1
 - If predictor outputs 1, then say the best expert outputs 0
 - Perfect expert is either in $[1, n/4]$, $[n/4+1, n/2]$, $[n/2+1, 3n/4]$, or $[3n/4+1, n]$
- ...
- Any deterministic strategy is incorrect on at least $\lg_2 n$ days

No Perfect Expert

- Suppose best expert makes M mistakes
- How can we guarantee we make at most $(M+1)(\log_2 n + 1)$ mistakes?
- Run Majority-and-Halving, but after throwing away all experts, bring them all back in and start over
- In each “phase”, each expert makes at least 1 mistake, and you make at most $\log_2 n + 1$ mistakes
- At most M finished phases, plus the last unfinished one

Doing Better

- If best expert makes M mistakes, we make at most $(M+1)(\log_2 n + 1) = O(M \log_2 n)$ mistakes
- Can't do better than best expert, who makes M mistakes
 - Suppose only one expert who always says 1 and is wrong M times
- Can't do better than $\log_2 n$ mistakes
- But can we make at most $\approx M + \log_2 n$ mistakes instead of $\approx M \cdot \log_2 n$?

Weighted Majority Algorithm

- Throwing away an expert when it makes a mistake is too drastic
- Assign weight w_i to i -th expert. Initialize all weights to 1
- On t -th day, compute sum of weights of experts who say 0, and sum of weights of experts who say 1
- Choose outcome with larger weight
- If an expert is wrong on day t , cut its weight in half

Weighted Majority Algorithm

- **Theorem:** If the best expert makes M mistakes, then the weighted majority algorithm makes at most $2.41(M + \log_2 n)$ mistakes!
- Proof: Let $\Phi = \sum_i w_i$. Initially $\Phi = n$
- When we make a mistake, $\Phi_{\text{new}} \leq \frac{3}{4} \cdot \Phi_{\text{old}}$
 - At least half of the weight (which made the majority prediction) gets halved (because it made a mistake)
- If we don't make a mistake, $\Phi_{\text{new}} \leq \Phi_{\text{old}}$

Weighted Majority Algorithm

- If we've made m mistakes so far, $\Phi_{\text{final}} \leq \left(\frac{3}{4}\right)^m \cdot \Phi_{\text{init}} = \left(\frac{3}{4}\right)^m \cdot n$
- Best expert i^* makes at most M mistakes, so $\Phi_{\text{final}} \geq w_{i^*} \geq \left(\frac{1}{2}\right)^M$
- So $\left(\frac{1}{2}\right)^M \leq \Phi_{\text{final}} \leq \left(\frac{3}{4}\right)^m \cdot n$, or $\left(\frac{4}{3}\right)^m \leq 2^M \cdot n$
- Taking logs, $m \leq \frac{M + \log_2 n}{\log_2 \left(\frac{4}{3}\right)} = 2.41(M + \log_2 n)$
- If best expert makes a mistake 10% of the time, we make a mistake 24% of the time, plus $((\log_2 n) / T) \%$, which is negligible with enough days

Improved Weighted Majority Algorithm

- **Only change:** if an expert is wrong on day t , multiply its weight by $1 - \epsilon$
- Still choose outcome given by the majority weight of experts in each day
- **Theorem:** If the best expert makes M mistakes, then the weighted majority algorithm makes at most $2(1 + \epsilon)M + O(\frac{\log_2 n}{\epsilon})$ mistakes

Improved Weighted Majority Algorithm

- Each time we make a mistake, $\Phi_{\text{new}} \leq (1 - \frac{\epsilon}{2}) \cdot \Phi_{\text{old}}$
 - At least half of the weight gets scaled by $1 - \epsilon$
- If we've made m mistakes, $\Phi_{\text{final}} \leq \left(1 - \frac{\epsilon}{2}\right)^m \cdot \Phi_{\text{init}} = \left(1 - \frac{\epsilon}{2}\right)^m n$
- $\Phi_{\text{final}} \geq w_{i^*} \geq (1 - \epsilon)^M$
- $(1 - \epsilon)^M \leq \Phi_{\text{final}} \leq \left(1 - \frac{\epsilon}{2}\right)^m n \text{ or } \frac{1}{\left(1 - \frac{\epsilon}{2}\right)^m} \leq \frac{n}{(1 - \epsilon)^M}$
- So $m \ln \frac{1}{1 - \frac{\epsilon}{2}} \leq M \cdot \ln \frac{1}{1 - \epsilon} + \ln n$
 - Use $\ln \frac{1}{1 - \frac{\epsilon}{2}} \geq \frac{\epsilon}{2}$ and $\ln \frac{1}{1 - \epsilon} \leq \epsilon + \epsilon^2$ for $\epsilon \in [0, \frac{1}{2}]$, and multiply both sides by $\frac{2}{\epsilon}$
 - $m \leq 2M(1 + \epsilon) + \frac{2 \ln n}{\epsilon}$

Lower Bound for Deterministic Algorithms

- **Theorem:** If the best expert makes M mistakes, then the weighted majority algorithm makes at most $2(1 + \epsilon)M + O(\frac{\log_2 n}{\epsilon})$ mistakes!
- If best expert is wrong 10% of the time, we're wrong 20% of the time
- $2M$ mistakes are necessary for any deterministic algorithm:
 - Suppose we have two experts - one always says 0 and one always says 1
 - If algorithm is deterministic, the adversary knows what prediction the algorithm will make on each day, so it can choose the opposite outcome
 - So algorithm incorrect on all days, but one expert is correct on at least half of the days

Randomized Weighted Majority Algorithm

- Assign weight w_i to i -th expert. Initialize all weights to 1
- On each day, predict 1 with probability $\frac{\sum_i \text{says 1} w_i}{\sum_i w_i}$, and predict 0 otherwise
- Equivalently, pick a random expert i with probability $\frac{w_i}{\sum_j w_j}$ and choose that expert's outcome
- When an expert makes a mistake, multiply its weight by $1 - \epsilon$

Randomized Weighted Majority Algorithm

- **Theorem:** If the best expert makes M mistakes, then the expected number of mistakes of the randomized weighted majority algorithm makes at most $(1 + \epsilon)M + \frac{\ln n}{\epsilon}$
- Previous $2M$ lower bound only applies to deterministic algorithms
- Let $\Phi = \sum_i w_i$. Initially $\Phi = n$
- Having fixed the outcome on all days, the potential varies deterministically

Randomized Weighted Majority Algorithm

- Let F_t be the fraction of total weight on the t -th day on experts that make a mistake on that day
- The expected number of mistakes we make is $\sum_t F_t$
- On day t : $\Phi_{\text{new}} = \Phi_{\text{old}} \cdot (1 - F_t) + \Phi_{\text{old}} \cdot F_t(1 - \epsilon) = \Phi_{\text{old}} (1 - \epsilon \cdot F_t)$
- $\Phi_{\text{final}} = n \cdot \prod_t (1 - \epsilon \cdot F_t) \leq n \cdot e^{-\epsilon \sum_t F_t}$ using $1 + x \leq e^x$ for all x
- Also, $\Phi_{\text{final}} \geq (1 - \epsilon)^M$

Randomized Weighted Majority Algorithm

- Have shown: $(1 - \epsilon)^M \leq \Phi_{\text{final}} \leq n \cdot e^{-\epsilon \sum_t F_t}$
- Taking natural logs, $\epsilon \sum_t F_t \leq M \ln \frac{1}{1-\epsilon} + \ln n$
- Using $\ln \frac{1}{1-\epsilon} \leq \epsilon + \epsilon^2$ for $\epsilon \in [0, \frac{1}{2}]$, and dividing both sides by ϵ we get:

$$\text{Expected number of mistakes} = \sum_t F_t \leq M(1 + \epsilon) + \frac{\ln n}{\epsilon}$$

Understanding the Error Rate

- Expected number of mistakes = $\sum_t F_t \leq M(1 + \epsilon) + \frac{\ln n}{\epsilon}$
- Best expert makes at most T mistakes, so $\sum_t F_t \leq M + \epsilon T + \frac{\ln n}{\epsilon}$
- Let M/T be optimal “error rate”
- Our expected error rate is at most optimal error rate + $\epsilon + \frac{\ln n}{\epsilon T}$
- Setting $\epsilon = \left(\frac{\ln n}{T}\right)^{1/2}$, our error rate \leq optimal rate + $2 \left(\frac{\ln n}{T}\right)^{1/2}$
- The last term is called the “regret”. As T gets larger, the regret goes to 0